ConvPE: Convolutional Position Embedding Improves Transformers

Wang Zehao
Nanjing University
wzhagi@163.com

ABSTRACT

ConvPE (Convolutional Position Embedding) substantially improves the performance and efficiency of Transformer models. With only 2.64% more parameters and a 2.71% increase in forward-pass time, it reduces training perplexity by 8.41%. Crucially, it achieves the baseline's final loss with 31.33% fewer training tokens, yielding a 26.87% net reduction in wall-clock training time. ConvPE also exhibits superior context extrapolation, reducing validation perplexity by 71.10% at 4K and 89.84% at 8K contexts. By generating windowed attention-score biases via a lightweight convolution, ConvPE remains compatible with KV-cache concatenation, various attention masks, and high-dimensional modalities. The code is publicly available at `https://github.com/wzhagi/convpe`.

Keywords Position Embedding · Language Models · Transformers · Attention · Vision Transformers

1 Introduction

Transformers [1] are insensitive to input order, whereas word order is crucial for language understanding. Explicit or implicit position embeddings are therefore necessary to inject order information into the model [2,3,4].

This paper proposes ConvPE, a simple and efficient position-embedding scheme that generates windowed attention-score biases with a lightweight convolution. This design allows ConvPE to switch freely between different attention masks, be combined with RoPE, support KV-cache concatenation, and naturally extend to high-dimensional modalities with n-dimensional convolution.

We comprehensively evaluate ConvPE in language-model pretraining and vision transformers. As summarized in Figure 1, ConvPE improves performance and efficiency across the board.

Performance:

ConvPE reduces training perplexity by 8.41% and improves average evaluation accuracy by 2.56%. It also transfers effectively to vision tasks, improving ViT-Tiny [5] accuracy on CIFAR-10 and CIFAR-100 [6] by 2.31% and 10.85%.

Training and Token Efficiency:

Despite a small 6.54% per-token wall-clock overhead and 2.64% more parameters, ConvPE reaches the RoPE baseline's performance with 31.33% fewer training tokens. This results in a net 26.87% reduction in total wall-clock training time to match the baseline.

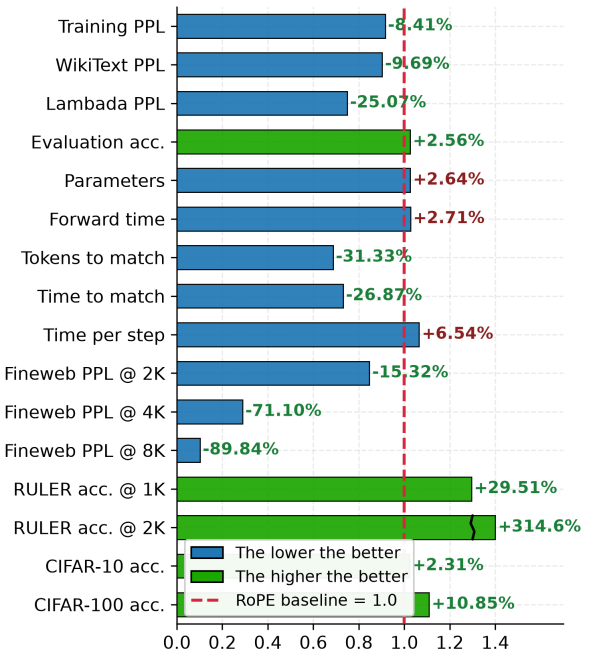


Figure 1: ConvPE / RoPE ratios.

Context Extension:

ConvPE significantly mitigates performance degradation at zero-shot context extensions. Compared to RoPE, it reduces validation perplexity by 15.32% at 2K, 71.10% at 4K, and 89.84% at 8K. On the RULER benchmark, ConvPE boosts accuracy by 31.67% at 1K and by over 300% at 2K and 4K contexts.

The main contributions of this paper are as follows:

- We propose ConvPE, a convolutional attention-score-bias position embedding that combines strong performance with broad compatibility.
- We study the influence of ConvPE hyperparameters and architecture choices, identifying a simple default configuration with a highly favorable performance-efficiency balance.
- We demonstrate that ConvPE significantly improves training efficiency, token efficiency, and zero-shot long-context extrapolation in language models.
- We verify the effectiveness of ConvPE2D in vision transformers on CIFAR-10 and CIFAR-100.
- We develop high-performance TileLang[7] kernels for ConvPE.

2 Background

2.1 Attention Mechanism

Given the input query $Q \in \mathbb{R}^{S_q \times D_k}$, key $K \in \mathbb{R}^{S_{kv} \times D_k}$, and value $V \in \mathbb{R}^{S_{kv} \times D_v}$, the computation of scaled dot-product attention[1] $O \in \mathbb{R}^{S_q \times D_v}$ is:

$$O = \text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}} + \mathbf{B}\right) V$$

S_q is the query sequence length. S_{kv} is the key/value sequence length. $\mathbf{B} \in \mathbb{R}^{S_q \times S_{kv}}$ is the attention score bias. Causal models introduce a causal mask $M \in \mathbb{R}^{S_q \times S_{kv}}$ that sets the attention score bias $\mathbf{B}$ for invisible positions to $-\infty$. With batch size and multi-head settings, Q, K, V, O additionally have B, H dimensions.

2.2 Position Embedding

Existing position embedding schemes can be roughly divided into five categories:

1. **Absolute Position Embedding**[1]: adds a fixed positional bias, usually sine waves, to the input vector. It is simple to implement, but its performance is usually limited.
2. **Learnable Positional Bias**[2,3]: adds a learnable positional bias to the input sequence. It is typically adopted in limited-window settings, such as spatial dimensions in vision models or fixed-rate compression.
3. **Rotary Position Embedding**[4]: encodes position as a rotation matrix and applies it to queries and keys, so that attention scores depend on relative positions. RoPE has been widely adopted in modern large language models, enabling long-context extension via frequency scaling, such as YaRN[8].
4. **Attention Score Bias**: adds a positional bias B to attention scores[9,10,11]. RoPE can be viewed as a rotation-based equivalent, as it also makes attention scores depend on relative positions.
5. **No Position Embedding (NoPE)**[12]: adds no explicit position embeddings; instead, it relies on model structure and causal attention masks to implicitly learn positional information requiring no explicit Position Embedding[13].

3 Method

Given token representations $X \in \mathbb{R}^{B \times S \times D}$, ConvPE first projects them to a small channel dimension $X_{down} \in \mathbb{R}^{B \times S \times C}$, then uses a convolution to generate windowed attention score bias $\mathbf{B}_{ConvPE} \in \mathbb{R}^{B \times S \times H_b \times W}$ for self-attention keys within the convolution window as shown in Figure 2.

$$\mathbf{B}^{b,h,s_q,s_{kv}} = \begin{cases} \mathbf{B}_{\text{ConvPE}}^{b,s_q,h \bmod H_b,w}, & w \in [0, W-1] \text{ where } w = s_{kv} - s_q + W - 1 \\ 0, & \text{otherwise.} \end{cases}$$

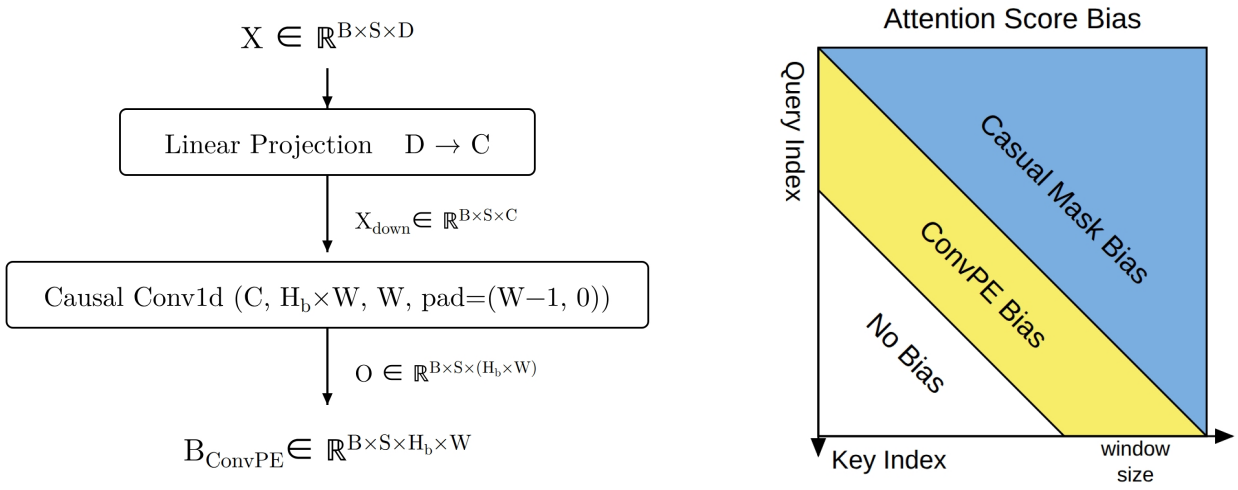


Figure 2: ConvPE Algorithm

The hyperparameters are:

- W: convolution window size.
- C: out features of linear projection and in channels of convolution.
- H_b: number of score bias heads.

The convolution takes C input channels, uses a kernel of size W, and outputs $H_b \times W$ channels. The projection has $D \times C$ parameters and the convolution has $W^2 \times C \times H_b$ parameters; the computation (MACs) per token is the same as the parameter count.

To reduce the convolution parameter count, we introduce three additional hyperparameters:

- Groups: number of convolution groups;
- Share: whether to share weights across convolution groups;
- Rank: rank used in the low-rank convolution.

4 Experiments

All models in our experiments use Multi-Query Attention (MQA)[14], with $K = V$[15]. Each attention layer has $D_{model} \times D_{head} \times (2 \times N_{heads} + 1)$ parameters. Each feed-forward layer uses SwiGLU[16] and has $8 \times D^2_{model}$ parameters.

4.1 Hyperparameter Search

We train each model for one epoch on WikiText-103[17] and record the validation loss. Training uses the Hugging Face Trainer[18], AdamW[19], and the GPT-2 tokenizer[3]. The training hyperparameters are: batch size 4, sequence length 1024, learning rate 3×10^{-4} with cosine learning-rate scheduler, and seed 42.

We train two types of backbone models, a tiny backbone with $N_{layers} = 12$, $D_{model} = 192$, $N_{heads} = 3$, $D_{head} = 64$, and a base backbone with $N_{layers} = 12$, $D_{model} = 768$, $N_{heads} = 12$, $D_{head} = 128$. The model name format is "`backbone RoPE/NoPE` W C Rank H_b Groups Share". For example, "`tiny rope w128d6r128h3g1s`" denotes a tiny backbone using RoPE with $W = 128$, $C = 6$, Rank $= 128$, $H_b = 3$, Groups $= 1$, and Share weights. For simplicity, model names may omit the default configuration: `NoPE`, `r0`(no low-rank convolution), `h1` ($H_b = 1$), `g1`(Groups $= 1$), and `s0` (no weight sharing).

First, we train tiny models with hyperparameters around "`nope w128d6r0h3g1`". Figure 3 shows that all ConvPE variants outperform the RoPE baseline. The largest gains come from increasing the window size and the projection dimension. Using grouped convolution causes the largest performance drop. RoPE integration and other parameter-saving methods show minor differences from "`nope w128d6r0h3g1`".

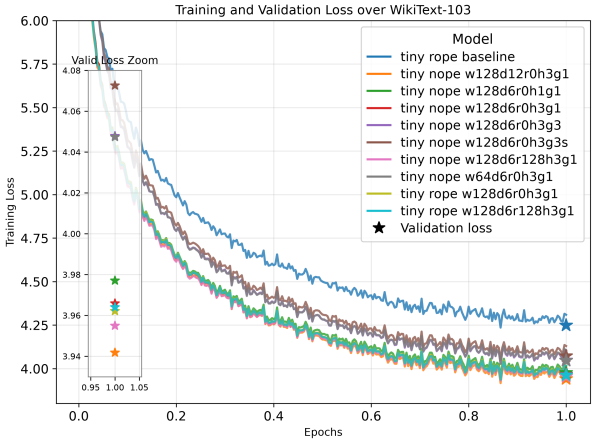


Figure 3: Tiny model training curve on WikiText-103

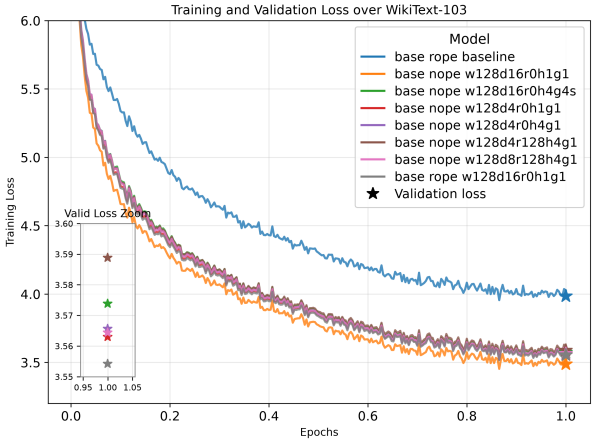


Figure 4: Base model training curve on WikiText-103

We then train base models with hyperparameters around "`nope w128d4r0h4g1`" and "`nope w128d16r0h1g1`". Figure 4 shows that $H_b = 1$ without any parameter-saving technique achieves the best balance between performance and parameter efficiency.

In additional experiments, we also try using X, Q[11] or KV as the convolution input, but performance drops.

4.2 Pretraining

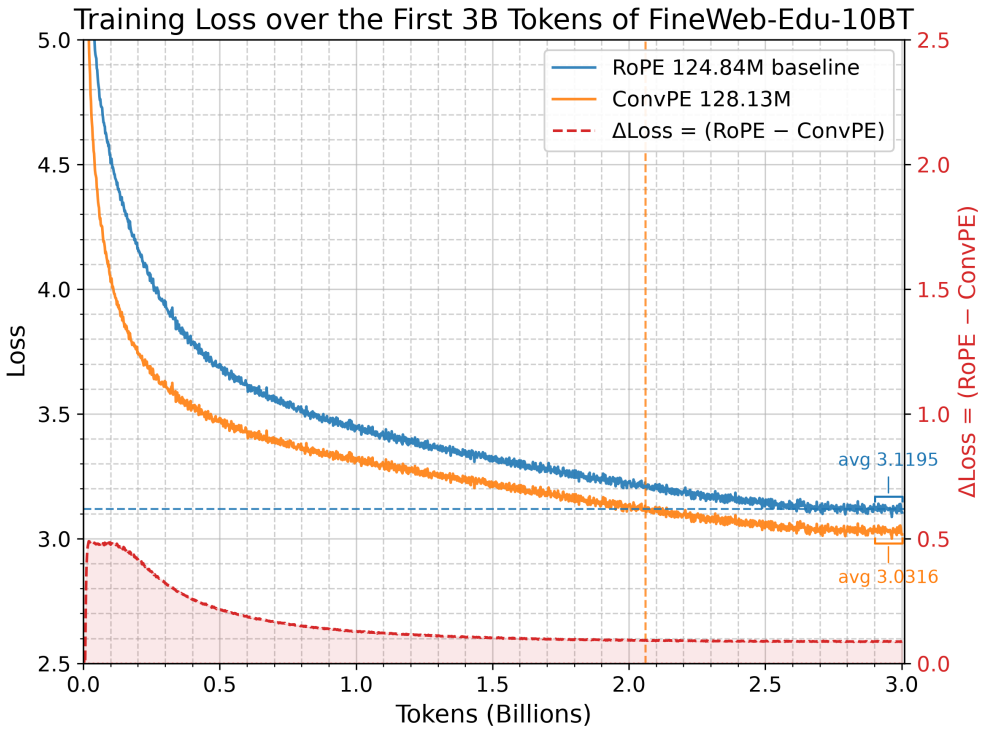


Figure 5: Pretraining curve on the first 3B tokens of FineWeb-Edu-10BT

We choose "`base w128d16`" as the ConvPE variant configuration for performance and throughput considerations. We train the RoPE baseline and the ConvPE variant from scratch on the first 3B tokens of FineWeb-Edu-10BT. We use Muon[20] to optimize the 2D parameters, excluding the embedding layer and the language head, with a learning rate of 3×10^{-3}, while Adam[21] optimizes the remaining parameters at a learning rate of 3×10^{-4}, both using a cosine

learning-rate scheduler. The training hyperparameters are: batch size 4, gradient accumulation steps 4, sequence length 1024, and seed 42.

Figure 5 shows that ConvPE converges faster and achieves a lower final loss compared to the RoPE baseline. We implement high-performance Flash MQA kernels in TileLang and benchmark two configurations: `base RoPE` with 124.84M parameters and `base w128d16` with 128.13M parameters ($+2.64\%$), using a batch size of 4 and a sequence length of 1024. While the forward-pass time increases by 2.71%, slightly exceeding the parameter increase, the final training loss drops from 3.1195 to 3.0316 (-2.82%). Despite the 6.54% per-step overhead, the number of training tokens required to match the RoPE baseline's final loss decreases from 3B to 2.06B, a 31.33% reduction. This translates to a net reduction in total wall-clock training time of approximately 26.87%.

4.3 Attention Analysis

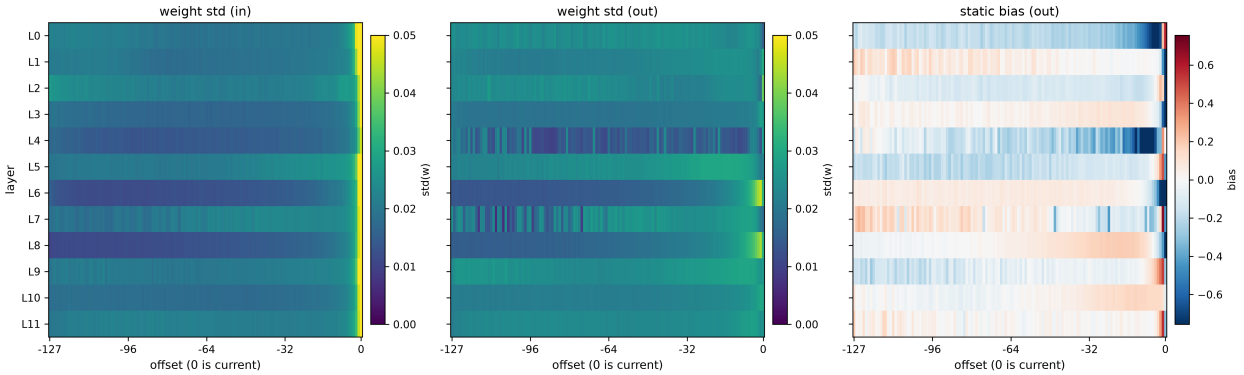


Figure 6: Visualization of ConvPE statistics across layers. Left: weight standard deviation in the input dimension. Middle: weight standard deviation in the output dimension. Right: static output bias. The x-axis represents relative offsets (-127 to 0), where 0 is the current token.

To understand the actual role of the ConvPE module, we analyze the weight standard deviation and the static output bias. Since the average weights are around zero, the standard deviation serves as an indicator of **feature selectivity**: a higher std implies the kernel focuses on specific channels. The numerical analysis reveals three primary properties:

1. Near-End Dominance. Near the current token, both the static bias and the input feature selectivity reach their maximum magnitudes. This indicates that the positional relationship between nearby tokens is crucial. ConvPE reflects this by applying highly specific feature extraction and strong positional priors to the immediate local context.

2. Far-End Relaxation. Towards the far end of the window, the standard deviation becomes flat, and the static bias converges to nearly zero. This shows that ConvPE does not impose strong static biases at this range, delegating the routing decision entirely to the convolution on dynamic input and the QK semantic matching of the tokens themselves.

3. Self-Suppression and Self-Enhancement. The static bias exhibits a clear layer-wise prior at the current position (offset 0). Layer 0 shows "self-enhancement" with a positive bias of $+0.5885$, whereas all deeper layers exhibit "self-suppression" with negative biases. This mechanism shifts attention in deeper layers away from the current token and toward the historical context.

4.4 Evaluations

This section evaluates the pretrained ConvPE and the RoPE baseline on pretraining and context-extension benchmarks without any fine-tuning.

4.4.1 Pretraining Evaluations

Table 1 shows that ConvPE consistently improves language-modeling performance over the RoPE baseline. On WikiText, perplexity drops from 46.56 to 42.05, and on LAMBADA perplexity decreases substantially from 118.10 to 88.49, with accuracy also improving from 0.2507 to 0.2604. ConvPE also outperforms RoPE on most downstream benchmarks. The gains are consistent across most metrics. The main exception is TruthfulQA-MC1, where both models fall below the random-choice baseline, indicating that this benchmark remains difficult for both models. Overall, the pretraining results support ConvPE as a stronger positional encoding than RoPE for the evaluated language tasks.

Dataset	Metric	RoPE	ConvPE	Random
wikitext	perplexity	46.56	**42.05**	-
lambada	perplexity, acc	118.10, 0.2507	**88.49**, **0.2604**	-
hellaswag	acc, acc norm	0.2894, 0.3148	**0.2949**, **0.3272**	0.25
xstorycloze_en	acc	0.5460	**0.5526**	0.5
piqa	acc, acc norm	0.6251, 0.6208	**0.6338**, **0.6235**	0.5
arc_easy	acc, acc norm	0.5114, 0.4487	**0.5421**, **0.4701**	0.25
winogrande	acc	0.5091	**0.5193**	0.5
boolq	acc	0.5676	**0.5893**	0.5
truthfulqa_mc1	acc	0.2399	0.2375	**0.25**

Table 1: Evaluation results of pretraining models

4.4.2 Context Extension Evaluations

Context	RoPE	ConvPE
1K	3.1247	**3.0312**
2K	3.1684	**3.0021**
4K	4.2690	**3.0278**
8K	5.6443	**3.3575**

Table 2: Validation loss on 1M training tokens

We first evaluate language-modeling performance using 1M tokens from the training dataset with increasing context length. Both models are pretrained on 1K-token contexts. Table 2 shows a clear advantage for ConvPE in language modeling as context length increases. RoPE degrades sharply beyond 2K, whereas ConvPE remains comparatively stable up to 4K and degrades much more gracefully at 8K. This indicates that ConvPE provides better context-length extrapolation and preserves model performance under longer contexts.

	1K		2K		4K	
Task	RoPE	ConvPE	RoPE	ConvPE	RoPE	ConvPE
niah_single_1	0.614	**0.890**	0.004	**0.722**	0.000	**0.314**
niah_single_2	0.612	**0.664**	0.000	**0.038**	0.000	0.000
niah_single_3	0.090	**0.180**	0.000	**0.014**	0.000	0.000
niah_multikey_1	0.198	**0.200**	0.118	**0.194**	0.000	0.002
niah_multikey_2	0.002	0.000	–	–	–	–
niah_multikey_3	0.000	0.008	–	–	–	–
niah_multiquery	0.104	**0.186**	0.029	**0.179**	0.000	0.000
niah_multivalue	0.149	**0.187**	0.127	**0.186**	0.000	0.000
ruler_vt	**0.032**	0.008	0.000	**0.052**	0.000	**0.107**
ruler_cwe	0.002	0.002	–	–	–	–
ruler_fwe	0.001	**0.012**	–	–	–	–
ruler_qa_squad	0.148	**0.191**	0.078	**0.091**	0.017	**0.052**

Table 3: RULER evaluations

We then evaluate both models on RULER tasks. Table 3 shows that ConvPE generally improves over RoPE on several long-context retrieval and reasoning tasks, particularly at 1K and 2K contexts. ConvPE even performs well at 4K contexts on some tasks. Some tasks remain near zero for both models, indicating that these tasks are too difficult for 125M-parameter models. Although the language-modeling loss of ConvPE remains stable up to 4K, the downstream tasks show that zero-shot context extension is still far from solved.

4.5 2D Position Embedding

To produce 2D positional embeddings, we use 2D convolutional layers that take a 2D window input and produce a 2D window output. We use ViT-Tiny[5] models with Multi-Head Attention (MHA) and the configuration $N_{\text{layers}} = 12$,

$D_{\text{model}} = 192$, $N_{\text{heads}} = 3$, and $D_{\text{head}} = 64$. We train the models on CIFAR-10 and CIFAR-100[6] for 100 epochs with patch shape 4×4.

model	params	CIFAR-10	CIFAR-100
RoPE baseline	5.346M	0.910	0.627
ConvPE2D `w3d4`	5.359M	0.931	**0.695**
ConvPE2D `w3d8`	5.372M	**0.933**	0.683
ConvPE2D `w5d4`	5.402M	0.929	0.682
ConvNet	4.335M	**0.945**	**0.737**

Table 4: Validation accuracy on CIFAR-10 and CIFAR-100

Table 4 shows that ConvPE transfers effectively to vision transformers and outperforms the RoPE baseline. The influence of hyperparameters requires further study on larger models and datasets.

5 Conclusions and Future Work

ConvPE delivers substantial performance gains at a minor parameter and computational cost, providing a powerful alternative for positional embeddings. We adopt $H_b = 1$ without any parameter-saving technique as the default configuration, which achieves the best balance among performance, parameter efficiency, and computational cost. ConvPE is compatible with RoPE, various attention masks, and KV-cache concatenation, and it naturally supports multidimensional extensions.

Future work can explore the following directions:

- **Kernel Optimization.** Further optimize the backward kernel and develop additional kernels for a wider range of hardware.
- **Scaling behavior.** Investigate ConvPE with larger Transformers and datasets, and identify configuration patterns at scale.
- **Long-context modeling.** Evaluate ConvPE in long-context settings, where positional generalization is particularly important.
- **Architectural integration.** Combine ConvPE with Sliding Window Attention (SWA)[22], Compressed Sparse Attention (CSA)[15], and other Transformer variants.
- **Multidimensional modalities.** Extend ConvPE to high-dimensional modalities such as vision, speech, and multimodal data.

References

[1] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

[2] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL-HLT*, pages 4171–4186, 2019.

[3] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. *OpenAI Blog*, 1(8):9, 2019.

[4] Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. Roformer: Enhanced transformer with rotary position embedding. *arXiv preprint arXiv:2104.09864*, 2021.

[5] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, et al. An image is worth 16x16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations*, 2021.

[6] Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.

[7] Lei Wang, Yu Cheng, Yining Shi, Zhengju Tang, Zhiwen Mo, Wenhao Xie, Lingxiao Ma, Yuqing Xia, Jilong Xue, Fan Yang, and Zhi Yang. Tilelang: A composable tiled programming model for AI systems. *CoRR*, abs/2504.17577, 2025. doi:10.48550/ARXIV.2504.17577. URL `https://doi.org/10.48550/arXiv.2504.17577`.

[8] Bowen Peng, Jeffrey Quesnelle, Honglu Fan, and Enrico Shippole. YaRN: Efficient context window extension of large language models. In *The Twelfth International Conference on Learning Representations*, 2024. URL https://openreview.net/forum?id=wHBfxhZu1u.

[9] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020.

[10] Ofir Press, Noah A. Smith, and Mike Lewis. Train short, test long: Attention with linear biases enables input length extrapolation. *arXiv preprint arXiv:2108.12409*, 2021.

[11] Tyler A. Chang, Yifan Xu, Weijian Xu, and Zhuowen Tu. Convolutions and self-attention: Re-interpreting relative positions in pre-trained language models, 2021. URL https://arxiv.org/abs/2106.05505.

[12] Amirhossein Kazemnejad, Inkit Padhi, Karthikeyan Natesan Ramamurthy, Payel Das, and Siva Reddy. The impact of positional encoding on length generalization in transformers. *arXiv preprint arXiv:2305.19466*, 2023.

[13] Kimi Team. Kimi k3: Open frontier intelligence. *arXiv preprint arXiv:2607.24653*, 2026. URL https://arxiv.org/abs/2607.24653.

[14] Noam Shazeer. Fast transformer decoding: One write-head is all you need. *arXiv preprint arXiv:1911.02150*, 2019.

[15] DeepSeek-AI. Deepseek-v4: Towards highly efficient million-token context intelligence. *arXiv preprint arXiv:2606.19348*, 2026. URL https://arxiv.org/abs/2606.19348.

[16] Noam Shazeer. Glu variants improve transformer. *arXiv preprint arXiv:2002.05202*, 2020.

[17] Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. Pointer sentinel mixture models. *arXiv preprint arXiv:1609.07843*, 2016.

[18] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Remi Louf, Morgan Funtowicz, et al. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 38–45, 2020.

[19] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019.

[20] Keller Jordan. Muon: An optimizer for hidden layers in neural networks. https://kellerjordan.github.io/posts/muon/, 2024.

[21] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In Yoshua Bengio and Yann LeCun, editors, *3rd International Conference on Learning Representations, ICLR 2015, San Diego, CA, USA, May 7-9, 2015, Conference Track Proceedings*, 2015. URL http://arxiv.org/abs/1412.6980.

[22] Iz Beltagy, Matthew E. Peters, and Arman Cohan. Longformer: The long-document transformer. *arXiv preprint arXiv:2004.05150*, 2020.